

Labeling AI-Generated Content May Not Change Its Persuasiveness

Isabel O. Gallegos, Chen Shani, Weiyan Shi, Federico Bianchi, Izzy Gainsburg, Dan Jurafsky, and Robb Willer

GENERATIVE AI TOOLS CAN NOW PRODUCE PERSUASIVE CONTENT AT PREVIOUSLY UNPRECEDENTED SCALE AND SPEED. There are many ways in which these tools can be used for positive impact in the world. But the emergence of persuasive, AI-generated content also makes possible many negative uses — such as influence operations, misinformation campaigns, and other kinds of deception — particularly in political contexts. These risks are compounded by a key issue: People struggle to distinguish AI-generated content from content written by humans, which helps influence campaigns and misinformation thrive.

These risks have led policymakers to call for authorship labels on AI-made content. In the European Union, for instance, the AI Act requires that entities deploying AI-generated or -manipulated content label it as such. In the United States, the AI Labeling Act and the AI Disclosure Act, both introduced in Congress in 2023 but not passed, would have implemented similar rules. Policymakers' calls for labels on AI-made content raise a key question: Will a label change how much the content influences people's political and public policy views?

Key Takeaways

In response to the rapidly improving ability of AI tools to create persuasive content, policymakers are increasingly calling for labels on AI-generated content—but little research has measured whether adding a label impacts the persuasiveness of the underlying messages.

We surveyed how more than 1,500 people perceive AI-generated policy messages when told the content had been created by an expert AI model, a human policy expert, or told nothing about its authorship.

Adding the label changed people's perceptions of whether the author was AI or human but did not significantly change the persuasiveness of the content itself, regardless of the policy domain (e.g., allowing colleges to pay student athletes) or participant demographics (e.g., political party).

Policy proposals requiring AI content labels may enhance transparency, but their inability to affect persuasiveness highlights the need for complementary safeguards (e.g., media literacy education) and ongoing research into how AI disclosure policies shape the information ecosystem.

In our paper “Labeling Messages as AI-Generated Does Not Reduce Their Persuasive Effects,” we surveyed a diverse group of Americans to investigate how adding authorship labels did or did not affect how they perceive AI-written policy appeals. Across four less-polarized public policy topics, we found no significant difference between people’s support for a policy argument when told the argument had been generated by an expert AI model, a human policy expert, or told nothing about its authorship. The labels also had no significant effects on people’s judgments of the content’s accuracy or people’s intentions to share the policy argument with others.

Policymakers’ calls for labels on AI-made content raise a key question: Will a label change how much the content influences people’s political and public policy views?

Policymakers should continue studying and debating effective AI disclosure policies, including how AI content labels may empower users to make more informed decisions about content consumption. Yet, while labels may enhance transparency, our work suggests that on their own, they may be insufficient in addressing the challenges posed by AI-generated information. Policymakers need to conduct further research and explore alternative interventions, including media literacy education and deamplification of AI-generated content.

Introduction

Previous work has explored the persuasiveness of AI-generated content without AI labels; perceptions of the credibility, reliability, or quality of labeled information; and the effect of different labels for AI-generated images on viewers’ beliefs. There remains a gap when it comes to research on the impact of attaching labels to AI-generated text on the content’s persuasiveness. We therefore chose to investigate how AI labels affect persuasiveness of messages on public policy issues.

There is good reason to assume that AI labels make people more skeptical of the underlying content. For instance, prior research has found that people generally prefer human content over AI content in news, public health messaging, donation solicitation, and social media content because they perceive the human content to be more trustworthy. On the other hand, AI labels could trigger people’s perceptions of technological expertise and sophistication, making them trust the AI-generated information more than if it were written by a human.

Our study tested both these hypotheses through a randomized experiment with 1,500 U.S. participants. We compared participants’ responses to AI-generated policy messages when labeled as AI-created versus human-authored versus unlabeled across four public policy domains: geoengineering, drug importation, college athlete salaries, and social media platform liability. We selected these four domains from the persuasion dataset as they are neither widely

discussed nor highly polarizing policy issues — and thus less likely to clash with people’s well-established or deeply held views.

After surveying participants’ prior knowledge about, support for, and confidence in their beliefs about one of the four public policy topics, we showed them an AI-generated argument about that policy area. We used OpenAI’s GPT-4o to generate the policy messages, manually editing the text only to correct factual errors. Examples of article titles include:

- *“Geoengineering poses too many risks and should not be considered.”*
- *“Drug importation jeopardizes safety controls and the domestic pharma industry.”*
- *“College athletes should be paid salaries.”*

All participants were shown the same message related to their assigned policy issue area. But they were variably told, at random, either that the message had been written by an expert AI model or a human policy expert, or they were given no details regarding authorship. Participants then indicated their levels of support for the policies, confidence in their support, their judgments of how accurate the message was, and their intentions to share it with others. They were also asked about their perceptions of the message’s source. We collected demographic data (e.g., political party affiliation, education level, age) so we could test how background characteristics impact the effects of AI labels.

*There remains a gap
when it comes to research on
the impact of attaching labels
to AI-generated text on the
content’s persuasiveness.*

Research Outcomes

We found that adding labels changes people’s perceptions of whether the content was authored by AI or a human — suggesting that labels improve transparency — but there was no evidence that these labels significantly change the *persuasiveness* of the content.

Of the participants who were told the message had been AI-written, 94.6% believed it, and of those told the message had been human-written, 89.3% believed it. Of those shown no label, 39% thought it was human-written, 31% thought it was AI-generated, and the remaining 30% were unsure.

The policy-related messages themselves were persuasive, in that they increased people’s support for the policy as compared to their baseline support, but participants’ support for policies did not significantly differ when messages were labeled as human- versus AI-made — or when there was no label at all. This trend persisted for the other three variables we measured: Across people’s confidence in their policy

We found that adding labels changes people’s perceptions of whether the content was authored by AI or a human — suggesting that labels improve transparency — but there was no evidence that these labels significantly change the persuasiveness of the content.

support, judgments of how accurate the message was, and intentions to share it with others, there was also no meaningful difference between the label conditions. These results may be attributable to the type of content we showed participants and the fact that we indicated that the content had been generated by an expert AI model. Because the advocacy messages were evidence-based and logical, and built on expert-based persuasion techniques, the chances that the information will be viewed skeptically are inherently lower.

Even for different subgroups of participants, our labels did not significantly impact the persuasiveness of political messages. We tested whether political party identity, prior knowledge about the topic, prior experience with AI tools, education level, and age influence the labels’ impact on policy support. While participants on the whole were not affected by AI labels, we did find that older individuals were less persuaded by content labeled as being AI-generated compared to content labeled as having been authored by a human. Otherwise, though, the labels had no effects in any subgroups of interest.

Policy Discussion

The rapid proliferation of AI-made content continues to provoke important questions about how we should mitigate negative uses and whether we must provide consumers of such content with more transparency — including through authorship labels. However, policy proposals that rely on AI labels alone to address the challenges of AI-generated information may be insufficient.

Based on our findings, telling a reader that text is AI-composed is unlikely to substantially affect its persuasiveness — meaning, for example, that placing an “AI-authored” label on a misleading but well-written policy argument is unlikely to mitigate the message’s harm. However, it is important to note that our study did not provide participants with variable information quality and veracity, which means we did not measure how labels reduce persuasiveness for one individual viewing different policy arguments. Testing AI labels with other kinds of information and in field settings may lead to different results. Further research is needed to better understand how AI disclosure policies shape trust and interactions within the information ecosystem, so that policymakers can make more informed decisions about possible interventions and their impacts.

Policymakers should also consider that our findings are influenced by the current state of LLM capabilities and their public perceptions. In the future, the style and ability of AI-generated persuasive communication could shift, as well as pervasive narratives around the trustworthiness of AI, potentially making AI labels more or less persuasive for readers.

Going forward, the impact of AI labels on persuasiveness will likely change as regulation that mandates disclosure develops — alongside people’s acceptance of or aversion to AI in general. These potential shifts, coupled with our study’s limitations — such as focusing on Americans, considering only two types of content labels, and not accounting for the other heuristics people may use to interpret information (beyond its source) — underscore the need for additional research as policymakers continue to debate AI content labels. Crucially, these issues also highlight the importance of implementing additional interventions to curb the spread of AI-generated material, such as media literacy education, which aims to equip people with the tools to discern AI-generated writing from human-authored writing as well as the validity of the information itself.

Policymakers should continue to consider the potential positive impacts of AI content label regulations and standards. These include bringing more transparency to online content, helping social media users make informed decisions about content consumption, and curbing the spread of AI-generated material in general. But if the policy goal is to decrease the persuasive effect of AI-generated messages, labels may not be the complete solution.

Further research is needed to better understand how AI disclosure policies shape trust and interactions within the information ecosystem.

Reference: The original article is accessible at Isabel O. Gallegos, Chen Shani et al., “**Labeling Messages as AI-Generated Does Not Reduce Their Persuasive Effects**,” arxiv.org, April 14, 2025, <https://arxiv.org/abs/2504.09865>.

[Stanford University’s Institute for Human-Centered Artificial Intelligence \(HAI\)](#) applies rigorous analysis and research to pressing policy questions on artificial intelligence. A pillar of HAI is to inform policymakers, industry leaders, and civil society by disseminating scholarship to a wide audience. HAI is a nonpartisan research institute, representing a range of voices. The views expressed in this policy brief reflect the views of the authors. For further information, please contact HAI-Policy@stanford.edu.



[Isabel O. Gallegos](#) is a PhD student in computer science at the Stanford School of Engineering and a JD candidate at Stanford Law School.



[Chen Shani](#) is a postdoctoral researcher at Stanford University.



[Weiyan Shi](#) is an assistant professor at Northeastern University, jointly appointed by the College of Engineering and Khoury College of Computer Sciences.



[Federico Bianchi](#) is a senior ML scientist at TogetherAI.



[Izzy Gainsburg](#) is a research scholar at Stanford University’s Center on Philanthropy and Civil Society.



[Dan Jurafsky](#) is the Reynolds Professor in Humanities, a professor of linguistics, and a professor of computer science at Stanford University.



[Robb Willer](#) is a professor of sociology, psychology (by courtesy), and organizational behavior at the Graduate School of Business (by courtesy) at Stanford University, director of the Politics and Social Change Lab, and co-director of Stanford’s Center on Philanthropy and Civil Society.



Stanford University
Human-Centered
Artificial Intelligence

Stanford HAI: 353 Jane Stanford Way, Stanford CA 94305-5008

T 650.725.4537 **F** 650.123.4567 **E** HAI-Policy@stanford.edu hai.stanford.edu