

Roundtable Report

Government AI Hire, Use, Buy (HUB) Roundtable Series

Roundtable 3: Government as a Buyer of AI

Authors

Carolina Oxenstierna

Aaron Snow

Danny Hague

Table of Contents

Introduction.....	2
Overview of Discussion.....	4
Key Themes	5
Flexibility and Adaptability	5
Risk Evaluation and Management.....	6
Agency Responsibilities	7
Liability	7
AI Inventory.....	7
Evaluating the Need for AI.....	8
Conclusion.....	9
Authors and Acknowledgments	10
Endnotes.....	11

Introduction

The U.S. government has embraced the potential of artificial intelligence. However, there are not yet sufficiently comprehensive strategies and guardrails across federal agencies to maximize AI applications that serve the public good while mitigating risks. In light of this gap, Georgetown's [Government AI Hire, Use, Buy \(HUB\) Roundtable Series](#) brings together leading voices to grapple with the legal liability questions that AI poses, examine AI's potential to transform government services, and consider how the government can better attract and use AI talent.

The federal government has been particularly active on AI over the last several years. Following the release of several advanced large language models, President Joe Biden signed the groundbreaking Executive Order on the Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence in October 2023.¹ In March 2024, the Office of Management and Budget released its government-wide policy, Advancing Governance, Innovation, and Risk Management for Agency Use of Artificial Intelligence, which complements the recently finalized European Union AI Act.² Additionally, the Government Accountability Office conducted its audit of current AI uses by federal agencies.³ This work followed efforts dating back to the first Trump administration, which implemented agency AI use case inventories in Executive Order 13960, Promoting the Use of Trustworthy Artificial Intelligence in the Federal Government.⁴ However, many questions on the U.S. government's use and procurement of AI tools and hiring of AI-enabling talent remain unaddressed.

Georgetown University's [Center for Security and Emerging Technology](#) and [Beeck Center for Social Impact and Innovation](#) (Beeck Center), together with the [Georgetown Law Institute for Technology Law and Policy](#), led a series of invite-only roundtables over the course of 2024 to discuss these pressing questions. **This resulting report was authored in 2024 after those discussions.**

On October 9, 2024, the Beeck Center hosted the third roundtable in the series. This event focused on the U.S. government's procurement process for AI. This report summarizes and reflects on the key findings, takeaways, and recommendations from the roundtable, which was conducted under the Chatham House Rule. All names, affiliations, and other identifying information have been omitted from this report.

The first two roundtables in the series featured discussions on the role of government as a user of AI and an employer of AI talent. These events featured debates on appropriately scoping definitions of AI to capture and govern potential AI use cases, establishing flexible regulatory frameworks that can effectively evolve with new

technology, and ensuring agencies have the projects and support systems in place to effectively onboard AI talent and incentivize them to stay in government over the long term.

Overview of Discussion

The roundtable included three sessions:

1. Foundations, transparency, and accountability.
2. Risk assessment, testing, and evaluation standards.
3. Security, liability, and legal frameworks.

The conversation covered a number of relevant issues, but a few key themes and recommendations emerged:

1. Agencies should establish flexible governance and risk management frameworks for acquiring AI systems.
2. Agencies have assumed a variety of core responsibilities that will need resourcing.
3. Policymakers and researchers should prioritize the development of policy and technical solutions, as well as resources, that address critical challenges in safety and security.

Key Themes

Flexibility and Adaptability

A key theme that emerged during the first session and echoed throughout the day was a need for flexible and adaptable approaches to governing agency AI acquisition. The first session focused on transparency, with significant time devoted to a discussion on what vendors ought to disclose to agencies when providing AI tools.

In addressing vendor data sharing with agencies, there was general agreement that certain expectations should be established up front. Some observed parallels to privacy-by-design debates and advocated for more protections than today's software default, which is to "collect all the data and figure out what to do with it later."

Several participants argued that having clear guidelines can help agencies and vendors, as there are many administrative and legal frameworks that require cross-disciplinary coordination. One participant argued that contracting professionals—who are experts in acquisition but not necessarily in technology—would benefit from guidance on the right questions to ask vendors to solicit the data and information agencies need. Another noted that vendors want to understand the "cost of compliance" as much as possible, so they can accurately forecast their investment in a new line of work. The participant noted that a lack of such clarity can stifle innovation, at least in the services available to government.

However, the group recognized that flexibility would ultimately be needed. Government access to vendor information is a negotiation, and striking the right balance is a back-and-forth process that would require additional resources to navigate. The varied nature of agency authorities—as well as differences in commercial off-the-shelf and non-COTS procurement processes—will dictate the need for flexible transparency requirements that are appropriately tailored to the specific use cases in question.

Participants also made note of government transparency to the public. They observed that transparency allows the public to hold the government accountable by ensuring citizens are aware of how they interact with AI when engaging with government services. Some argued that such transparency should include consistently updating use case inventories, documenting failures of AI systems, and establishing mechanisms for due diligence, such as investigations and reporting—all critical steps that allow civil society and other experts to track and evaluate government acquisition and use of AI.

Risk Evaluation and Management

Much of the day focused on how developers and agencies can effectively evaluate and manage risks from AI systems deployed in government. One participant emphasized that the risk of AI use in government comes from its deployment to the public, not the AI system itself. However, some participants noted that new risks introduced by generative AI, such as cybersecurity vulnerabilities in AI-generated code, may materialize in similar ways across multiple contexts and may be best mitigated by developers.

Some participants highlighted the existing legal frameworks that can be applied to governing and managing risks. While there was no consensus on which framework should be followed, participants did largely agree that any future approach should require more extensive governance and risk management practices for sensitive AI deployments. They specifically referenced how the advent of cybersecurity in the government could have been better managed if existing legal frameworks had been integrated quickly, using this as a cautionary lesson for AI risk management.

One participant highlighted the risk of model drift and argued for incorporating checks into the code that can flag when AI models start to drift. The participant called for continuous monitoring and testing, funding for audit teams, and integrated product teams that can ensure these protections are embedded into systems. The participant also advocated for training, once AI tools are in the door, to ensure government employees understand model drift and can detect and respond to it when it occurs.

Another participant shared what one private sector risk-evaluation process looks like. The approach involves regular review and risk categorization of products, with higher-risk systems receiving more intensive scrutiny from more senior officials.

Core to this process is regularly reviewing models and tools as they are developed, ensuring repeated risk evaluation along the way. This participant argued for specific triggers that ensure reevaluation is occurring, such as mandatory review when new features are added or the tool is adapted for a new mission or function.

The group also emphasized the need to consider existing risk management frameworks, noting the significant number of compliance plans already in place at agencies. Central to these comments was the need for reasonable, established checks and balances that ensure an iterative, responsive, and adaptable approach to evaluating and managing risk. These existing frameworks can help vendors and the agencies they serve navigate risks from greater use of AI in government.

Agency Responsibilities

A second key theme weaved throughout the day's sessions was the core responsibilities and liabilities that agencies carry in seeking to incorporate AI into their operations.

Liability

The crux of the day's third session centered on where liability sits when an AI system developed in the private sector and deployed by an agency falls short on safety or security. The prevailing sentiment in the room was that agencies bear the liability for the systems they deploy. To paraphrase one participant, agencies are the ones providing public services, and the liability and accountability to the public is theirs. This is something that participants argued agencies have already accepted as necessary to avoid overly narrowing the pool of available vendors. However, the group did discuss the benefits of considering a regime in which vendors could share some liability, if such a framework is possible.

AI Inventory

While liability was one of the more explicit ownership roles placed on agencies in the discussion, participants also called on agencies to do more to track their acquisition and use of AI.

One participant noted that the number of vendors selling AI systems to the government—and the lack of a universal system integrating these transactions—makes it difficult to catalog them. But the work is important in order to understand the ecosystem of AI in government and conduct the regular risk management and evaluation described above. Another participant pointed to an example of an agency buying redundant software licenses from the same company, underscoring the lack of rigor in current tracking.

While no one disagreed on the importance of this work, several argued that tracking AI acquisitions is easier said than done and emphasized the importance of more resourcing, particularly for administrative offices that are viewed as overhead and not central to the agency's core mission.

Evaluating the Need for AI

Finally, one participant called on agencies to do a more rigorous assessment when considering a potential acquisition as to whether AI is even needed to fulfill the objective. The participant observed that excitement, particularly among the technologists in government, has led many to pursue AI tools when simpler solutions are available to meet the same need. Several participants agreed and noted that cost-benefit analyses are needed as to whether the advantages of AI in a given context outweigh the privacy, safety, and security risks.

Conclusion

A central takeaway from the roundtable is that there are a number of unanswered policy and technical questions when it comes to AI. As a result, agencies will need to take ownership of developing tools and implementing rules that can help guide the responsible acquisition of AI. Equally important, policymakers will need to resource agencies for these efforts, funding programs such as privacy and compliance offices charged with managing risk and agency officials looking to catalog and document government use of AI.

Ultimately, these efforts will need to be iterative and flexible so that agencies can adapt according to new technological developments. We know these advances will come, but we do not know exactly what they will look like. An agency acquisition model that has safeguards in place, but flexibility to update as needed, will be key to unlocking the power of AI while limiting its potential for harm.

Authors


Carolina Oxenstierna is a senior at Georgetown University's Walsh School of Foreign Service, studying AI governance and cybersecurity.

Aaron Snow is a former fellow at Georgetown University's Beeck Center for Social Impact and Innovation.

Danny Hague is the assistant director of external affairs and strategy at Georgetown University's Center for Security and Emerging Technology.

Acknowledgments

This work was supported by the Rockefeller Foundation. The authors sincerely appreciate all of the effort and insight provided by roundtable participants. A specific thank-you is owed to all those who helped execute the roundtable and publication of this report, including Igor Mikolic-Torreira, Owen Daniels, and Shelton Fitch.



© 2025 by the Center for Security and Emerging Technology. This work is licensed under a Creative Commons Attribution-Non Commercial 4.0 International License.

To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc/4.0/>.

Document Identifier: doi: 10.51593/ 20240053

Endnotes

¹ Exec. Order No. 14110, 88 FR 75191 (2023), www.federalregister.gov/documents/2023/11/01/2023-24283/safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence.

² Director of OMB, *Advancing Governance, Innovation, and Risk Management for Agency Use of Artificial Intelligence*, M-24-10, March 28, 2024, www.whitehouse.gov/wp-content/uploads/2024/03/M-24-10-Advancing-Governance-Innovation-and-Risk-Management-for-Agency-Use-of-Artificial-Intelligence.pdf; European Union, “Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 Laying Down Harmonised Rules on Artificial Intelligence and Amending Regulations (EC) No 300/2008, (EU) No 167/2013, (EU) No 168/2013, (EU) 2018/858, (EU) 2018/1139 and (EU) 2019/2144 and Directives 2014/90/EU, (EU) 2016/797 and (EU) 2020/1828 (Artificial Intelligence Act)Text with EEA Relevance” (2024), <https://artificialintelligenceact.eu/the-act/>.

³ Government Accountability Office, *Artificial Intelligence: Agencies Have Begun Implementation but Need to Complete Key Requirements* (Washington, D.C.: Government Accountability Office, December 2023), www.gao.gov/assets/d24105980.pdf.

⁴ Exec. Order No. 13960, 85 FR 78939 (2020), www.federalregister.gov/documents/2020/12/08/2020-27065/promoting-the-use-of-trustworthy-artificial-intelligence-in-the-federal-government.